

## Importing the required packages

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, precision_recall_fscore
import joblib
import pickle
```

## Loading the dataset

```
In [2]: import os, types
import pandas as pd
from botocore.client import Config
import ibm_boto3

def __iter__(self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your cr
# You might want to remove those credentials before you share the notebook.
cos_client = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='PM77kN4FyWauT9b7yYQ1Mgxi4SbBI4_3cKxpFreNFXdY',
    ibm_auth_endpoint="https://iam.cloud.ibm.com/oidc/token",
    config=Config(signature_version='oauth'),
    endpoint_url='https://s3.private.us.cloud-object-storage.appdomain.cloud')

bucket = 'flightdelayprediction-donotdelete-pr-3uv7e82sbqe8r'
object_key = 'flightdata.csv'

body = cos_client.get_object(Bucket=bucket, Key=object_key)['Body']
# add missing __iter__ method, so pandas accepts body as file-like object
if not hasattr(body, "__iter__"): body.__iter__ = types.MethodType( __iter__, body )

df = pd.read_csv(body)
df.head()
```

```
Out[2]:
```

	YEAR	QUARTER	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	UNIQUE_CARRIER	TAIL_NUM	FL_NUM	C
0	2016	1	1	1	5	DL	N836DN	1399	
1	2016	1	1	1	5	DL	N964DN	1476	
2	2016	1	1	1	5	DL	N813DN	1597	
3	2016	1	1	1	5	DL	N587NW	1768	
4	2016	1	1	1	5	DL	N836DN	1823	

5 rows × 26 columns


```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11231 entries, 0 to 11230
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   YEAR                  11231 non-null  int64
1   QUARTER               11231 non-null  int64
2   MONTH                11231 non-null  int64
3   DAY_OF_MONTH          11231 non-null  int64
4   DAY_OF_WEEK           11231 non-null  int64
5   UNIQUE_CARRIER       11231 non-null  object
6   TAIL_NUM              11231 non-null  object
7   FL_NUM               11231 non-null  int64
8   ORIGIN_AIRPORT_ID     11231 non-null  int64
9   ORIGIN                11231 non-null  object
10  DEST_AIRPORT_ID       11231 non-null  int64
11  DEST                  11231 non-null  object
12  CRS_DEP_TIME          11231 non-null  int64
13  DEP_TIME              11124 non-null  float64
14  DEP_DELAY             11124 non-null  float64
15  DEP_DEL15             11124 non-null  float64
16  CRS_ARR_TIME          11231 non-null  int64
17  ARR_TIME              11116 non-null  float64
18  ARR_DELAY             11043 non-null  float64
19  ARR_DEL15             11043 non-null  float64
20  CANCELLED             11231 non-null  float64
21  DIVERTED              11231 non-null  float64
22  CRS_ELAPSED_TIME      11231 non-null  float64
23  ACTUAL_ELAPSED_TIME   11043 non-null  float64
24  DISTANCE              11231 non-null  float64
25  Unnamed: 25           0 non-null      float64
dtypes: float64(12), int64(10), object(4)
memory usage: 2.2+ MB
```

## Performing Univariate Analysis

### Using Pie Chart

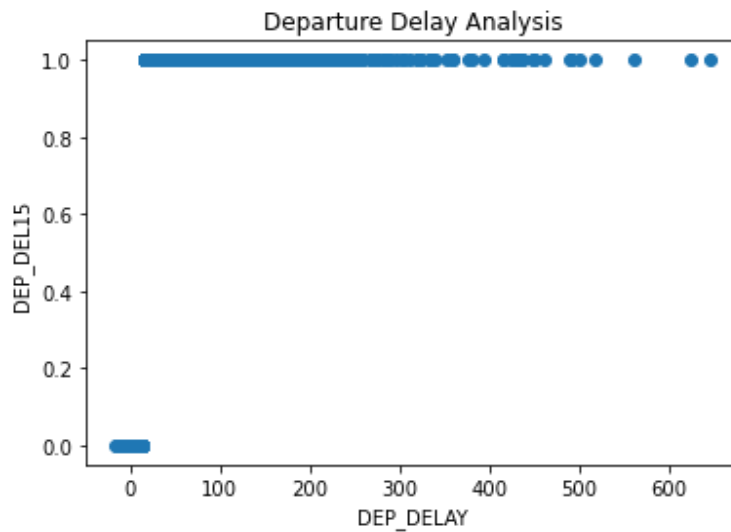
```
In [4]: df['YEAR'].value_counts().plot(kind='pie', autopct='%0.0f')
plt.show()
```



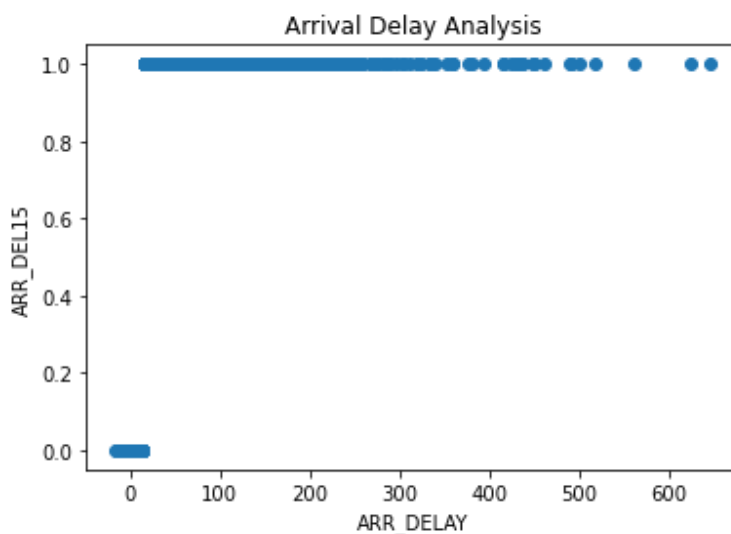
## Performing Bivariate Analysis

### Using scatterplot

```
In [5]: plt.scatter(df.DEP_DELAY, df.DEP_DEL15)
plt.title('Departure Delay Analysis')
plt.xlabel('DEP_DELAY')
plt.ylabel('DEP_DEL15')
plt.show()
```




```
In [6]: plt.scatter(df.DEP_DELAY, df.DEP_DEL15)
plt.title('Arrival Delay Analysis')
plt.xlabel('ARR_DELAY')
plt.ylabel('ARR_DEL15')
plt.show()
```




## Using lineplots


```
In [7]: fig, ax = plt.subplots(figsize=(10, 3))
plt.subplot(1, 2, 1)
plt.title('CRS_DEP_TIME')
plt.plot(df.CRS_DEP_TIME)
plt.subplot(1, 2, 2)
plt.title('DEP_TIME')
plt.plot(df.DEP_TIME)
plt.show()
```



```
In [8]: fig, ax = plt.subplots(figsize=(10, 3))
plt.subplot(1, 2, 1)
plt.title('CRS_ARR_TIME')
plt.plot(df.CRS_ARR_TIME)
plt.subplot(1, 2, 2)
plt.title('ARR_TIME')
plt.plot(df.ARR_TIME)
plt.show()
```



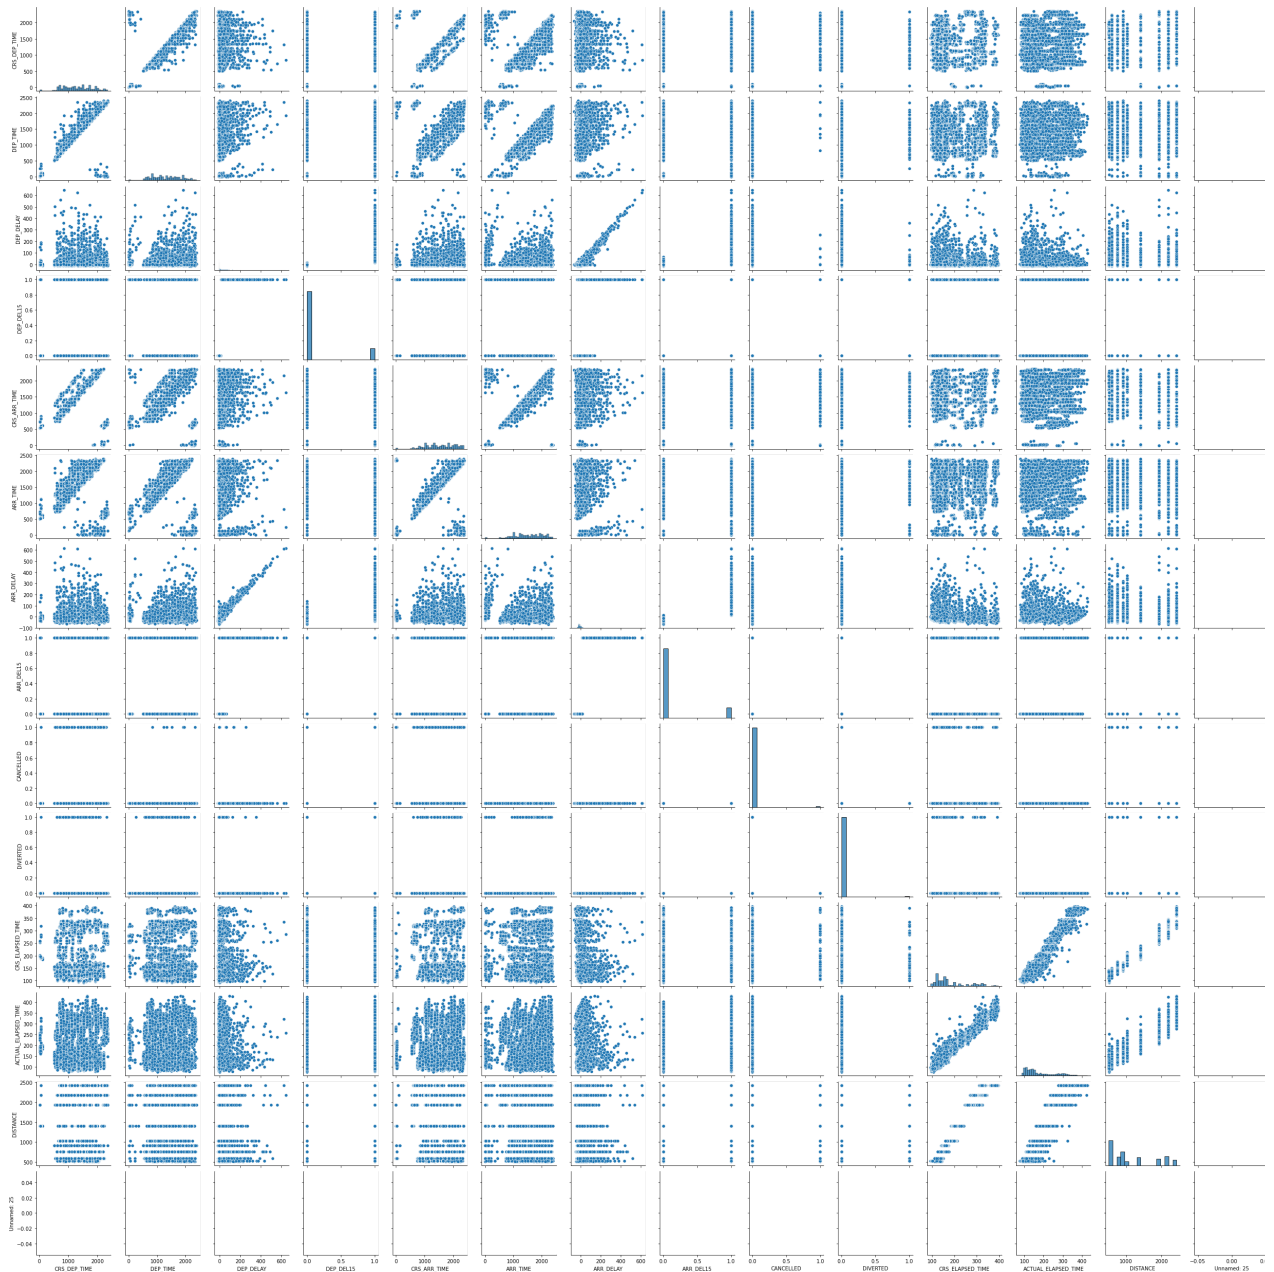
```
In [9]: fig, ax = plt.subplots(figsize=(10, 3))
plt.subplot(1, 2, 1)
plt.title('CRS_ELAPSED_TIME')
plt.plot(df.CRS_ELAPSED_TIME)
plt.subplot(1, 2, 2)
plt.title('ACTUAL_ELAPSED_TIME')
plt.plot(df.ACTUAL_ELAPSED_TIME)
plt.show()
```



## Performing Multivariate Analysis


### Using pairplot

```
In [10]: sb.pairplot(df.iloc[:, 12:])
plt.show()
```



## Using heatmap

```
In [11]: fig, ax = plt.subplots(figsize=(15, 10))
sb.heatmap(df.iloc[:, 12:].corr(), annot=True, ax=ax)
plt.show()
```



## Performing Descriptive Analysis


In [12]:

df.describe()

Out[12]:

	YEAR	QUARTER	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	FL_NUM	ORIGIN_AIRPC
count	11231.0	11231.000000	11231.000000	11231.000000	11231.000000	11231.000000	11231.000000
mean	2016.0	2.544475	6.628973	15.790758	3.960199	1334.325617	12334.5
std	0.0	1.090701	3.354678	8.782056	1.995257	811.875227	1595.0
min	2016.0	1.000000	1.000000	1.000000	1.000000	7.000000	10397.0
25%	2016.0	2.000000	4.000000	8.000000	2.000000	624.000000	10397.0
50%	2016.0	3.000000	7.000000	16.000000	4.000000	1267.000000	12478.0
75%	2016.0	3.000000	9.000000	23.000000	6.000000	2032.000000	13487.0
max	2016.0	4.000000	12.000000	31.000000	7.000000	2853.000000	14747.0

8 rows × 22 columns



## Dropping unnecessary columns

In [13]:

df = df[['FL\_NUM', 'MONTH', 'DAY\_OF\_MONTH', 'DAY\_OF\_WEEK', 'ORIGIN', 'DEST', 'DEP\_DEL15'], df.head()

Out[13]:

	FL_NUM	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	ORIGIN	DEST	DEP_DEL15	CRS_ARR_TIME	ARR_I
0	1399	1	1	5	ATL	SEA	0.0	2143	
1	1476	1	1	5	DTW	MSP	0.0	1435	
2	1597	1	1	5	ATL	SEA	0.0	1215	
3	1768	1	1	5	SEA	MSP	0.0	1335	
4	1823	1	1	5	SEA	DTW	0.0	607	

## Handling Missing Values

### Checking for null values

In [14]: `df.isnull().any()`

Out[14]:

FL_NUM	False
MONTH	False
DAY_OF_MONTH	False
DAY_OF_WEEK	False
ORIGIN	False
DEST	False
DEP_DEL15	True
CRS_ARR_TIME	False
ARR_DEL15	True

dtype: bool

### Replacing null values

In [15]:

```
df.fillna(df['DEP_DEL15'].mode()[0], inplace=True)
df.fillna(df['ARR_DEL15'].mode()[0], inplace=True)
```

### Checking if the replacement is made

In [16]: `df.isnull().any()`

Out[16]:


FL_NUM	False
MONTH	False
DAY_OF_MONTH	False
DAY_OF_WEEK	False
ORIGIN	False
DEST	False
DEP_DEL15	False
CRS_ARR_TIME	False
ARR_DEL15	False

dtype: bool

## Handling Outliers

In [17]:

```
fig, ax = plt.subplots(figsize=(5, 6))
sb.boxplot(data=df['CRS_ARR_TIME'])
plt.show()
```



There are no outliers

## Encoding

### One Hot Encoding

```
In [18]: df = pd.get_dummies(df, columns=['ORIGIN', 'DEST'])
df.head()
```

```
Out[18]:
```

	FL_NUM	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	DEP_DEL15	CRS_ARR_TIME	ARR_DEL15	ORIGIN_ATL	ORIGIN_DTW	ORIGIN_JFK	ORIGIN_MSP	ORIGIN_SEA	DEST_ATL	DEST_DTW	DEST_JFK	DEST_MSP	DEST_SEA
0	1399	1	1	5	0.0	2143	0.0										
1	1476	1	1	5	0.0	1435	0.0										
2	1597	1	1	5	0.0	1215	0.0										
3	1768	1	1	5	0.0	1335	0.0										
4	1823	1	1	5	0.0	607	0.0										

```
In [19]: df.columns
```

```
Out[19]: Index(['FL_NUM', 'MONTH', 'DAY_OF_MONTH', 'DAY_OF_WEEK', 'DEP_DEL15',
        'CRS_ARR_TIME', 'ARR_DEL15', 'ORIGIN_ATL', 'ORIGIN_DTW', 'ORIGIN_JFK',
        'ORIGIN_MSP', 'ORIGIN_SEA', 'DEST_ATL', 'DEST_DTW', 'DEST_JFK',
        'DEST_MSP', 'DEST_SEA'],
        dtype='object')
```

## Splitting dataset into Independent and Dependent Variables

```
In [20]: X = df.drop(columns=['ARR_DEL15'])
Y = df[['ARR_DEL15']]
```



## Scaling the Independent Variables

```
In [21]: scale_crs = StandardScaler()
X[['CRS_ARR_TIME']] = scale_crs.fit_transform(X[['CRS_ARR_TIME']])
scale_flnum = StandardScaler()
X[['FL_NUM']] = scale_flnum.fit_transform(X[['FL_NUM']])
X.head()
```

```
Out[21]:
```

	FL_NUM	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	DEP_DEL15	CRS_ARR_TIME	ORIGIN_ATL	ORIGIN_
0	0.079664	1	1	5	0.0	1.205371	1	
1	0.174510	1	1	5	0.0	-0.203612	0	
2	0.323555	1	1	5	0.0	-0.641431	1	
3	0.534188	1	1	5	0.0	-0.402620	0	
4	0.601935	1	1	5	0.0	-1.851405	0	

```
In [22]: X.FL_NUM.value_counts()
```

```
Out[22]:
```

-0.549771	98
-0.918071	96
0.808873	96
-0.919302	95
-0.532526	94
..	
1.865732	1
0.242258	1
0.195451	1
0.212695	1
1.608292	1

Name: FL\_NUM, Length: 690, dtype: int64

## Converting the Independent and Dependent Variables to 1D Arrays

```
In [23]: X = X.values
Y = Y.values
```

## Splitting dataset into Train and Test datasets

```
In [24]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2)
```

```
In [25]: X_train.shape, X_test.shape, Y_train.shape, Y_test.shape
```

```
Out[25]: ((8984, 16), (2247, 16), (8984, 1), (2247, 1))
```

## Building the Machine Learning Models

### Logistic Regression

```
In [26]: from sklearn.linear_model import LogisticRegression
log_reg = LogisticRegression(max_iter=200)
log_reg.fit(X_train, Y_train.ravel())
```

```
Out[26]: LogisticRegression(max_iter=200)
```

## Support Vector Machine (Classifier)

```
In [27]: from sklearn.svm import SVC  
svc = SVC()  
svc.fit(X_train, Y_train.ravel())
```

```
Out[27]: SVC()
```

## KNN Classifier

```
In [28]: from sklearn.neighbors import KNeighborsClassifier  
knn = KNeighborsClassifier()  
knn.fit(X_train, Y_train.ravel())
```

```
Out[28]: KNeighborsClassifier()
```

## Random Forest Classifier

```
In [29]: from sklearn.ensemble import RandomForestClassifier  
rf = RandomForestClassifier(n_estimators=15, max_depth=3)  
rf.fit(X_train, Y_train.ravel())
```

```
Out[29]: RandomForestClassifier(max_depth=3, n_estimators=15)
```

## Testing the Models

### Logistic Regression

```
In [30]: Y_pred_log_train = log_reg.predict(X_train)  
Y_pred_log_test = log_reg.predict(X_test)
```

```
In [31]: pd.DataFrame(Y_pred_log_train).value_counts()
```

```
Out[31]: 0.0    7692  
1.0    1292  
dtype: int64
```

```
In [32]: pd.DataFrame(Y_pred_log_test).value_counts()
```

```
Out[32]: 0.0    1960  
1.0     287  
dtype: int64
```

## Support Vector Machine (Classifier)

```
In [33]: Y_pred_svc_train = svc.predict(X_train)  
Y_pred_svc_test = svc.predict(X_test)
```

```
In [34]: pd.DataFrame(Y_pred_svc_train).value_counts()
```

```
Out[34]: 0.0    7685  
1.0    1299  
dtype: int64
```

```
In [35]: pd.DataFrame(Y_pred_svc_test).value_counts()
```

```
Out[35]: 0.0    1957
         1.0     290
         dtype: int64
```

## KNN Classifier

```
In [36]: Y_pred_knn_train = knn.predict(X_train)
         Y_pred_knn_test  = knn.predict(X_test)
```

```
In [37]: pd.DataFrame(Y_pred_knn_train).value_counts()
```

```
Out[37]: 0.0    8586
         1.0     398
         dtype: int64
```

```
In [38]: pd.DataFrame(Y_pred_knn_test).value_counts()
```

```
Out[38]: 0.0    2176
         1.0      71
         dtype: int64
```

## Random Forest Classifier

```
In [39]: Y_pred_rf_train = rf.predict(X_train)
         Y_pred_rf_test  = rf.predict(X_test)
```

```
In [40]: pd.DataFrame(Y_pred_rf_train).value_counts()
```

```
Out[40]: 0.0    8976
         1.0       8
         dtype: int64
```

```
In [41]: pd.DataFrame(Y_pred_rf_test).value_counts()
```

```
Out[41]: 0.0    2247
         dtype: int64
```

## Evaluating the ML Models using Metrics

### Logistic Regression

#### Classification Report

```
In [42]: print(classification_report(Y_test, Y_pred_log_test))
```

	precision	recall	f1-score	support
0.0	0.96	0.94	0.95	1997
1.0	0.61	0.70	0.65	250
accuracy			0.92	2247
macro avg	0.78	0.82	0.80	2247
weighted avg	0.92	0.92	0.92	2247

#### Accuracy, Precision, Recall, F1 Score

```
In [43]: acc_log = accuracy_score(Y_test, Y_pred_log_test)
         prec_log, rec_log, f1_log, sup_log = precision_recall_fscore_support(Y_test, Y_pred_log_test)
```

```
print('Accuracy Score =', acc_log)
print('Precision =', prec_log[0])
print('Recall =', rec_log[0])
print('F1 Score =', f1_log[0])
```

Accuracy Score = 0.9158878504672897  
 Precision = 0.9612244897959183  
 Recall = 0.943415122684026  
 F1 Score = 0.9522365428354814

### Checking for Overfitting and Underfitting

```
In [44]: log_train_acc = accuracy_score(Y_train, Y_pred_log_train)
log_test_acc = accuracy_score(Y_test, Y_pred_log_test)
print('Training Accuracy =', log_train_acc)
print('Testing Accuracy =', log_test_acc)
```

Training Accuracy = 0.9213045414069457  
 Testing Accuracy = 0.9158878504672897

**There is no big variation in the training and testing accuracy. Therefore, the Logistic Regression model is not overfit or underfit.**

### Confusion Matrix

```
In [45]: pd.crosstab(Y_test.ravel(), Y_pred_log_test)
```

```
Out[45]:
```

	col_0	0.0	1.0
row_0			
0.0	1884	113	
1.0	76	174	

## Support Vector Machine (Classifier)

### Classification Report

```
In [46]: print(classification_report(Y_test, Y_pred_svc_test))
```

	precision	recall	f1-score	support
0.0	0.96	0.94	0.95	1997
1.0	0.61	0.70	0.65	250
accuracy			0.92	2247
macro avg	0.78	0.82	0.80	2247
weighted avg	0.92	0.92	0.92	2247

### Accuracy, Precision, Recall, F1 Score

```
In [47]: acc_svc = accuracy_score(Y_test, Y_pred_svc_test)
prec_svc, rec_svc, f1_svc, sup_svc = precision_recall_fscore_support(Y_test, Y_pred_svc_test)
print('Accuracy Score =', acc_svc)
print('Precision =', prec_svc[0])
print('Recall =', rec_svc[0])
print('F1 Score =', f1_svc[0])
```

Accuracy Score = 0.9163328882955051  
 Precision = 0.9621870209504343  
 Recall = 0.942914371557336  
 F1 Score = 0.9524532119372786

### Checking for Overfitting and Underfitting

```
In [48]: svc_train_acc = accuracy_score(Y_train, Y_pred_svc_train)
svc_test_acc = accuracy_score(Y_test, Y_pred_svc_test)
print('Training Accuracy =', svc_train_acc)
print('Testing Accuracy =', svc_test_acc)
```

Training Accuracy = 0.9214158504007124  
 Testing Accuracy = 0.9163328882955051

**There is no big variation in the training and testing accuracy. Therefore, the Support Vector Classifier model is not overfit or underfit.**

### Confusion Matrix

```
In [49]: pd.crosstab(Y_test.ravel(), Y_pred_svc_test)
```

```
Out[49]:   col_0   0.0   1.0
row_0
0.0   1883   114
1.0     74   176
```

## KNN Classifier

### Classification Report

```
In [50]: print(classification_report(Y_test, Y_pred_knn_test))
```

	precision	recall	f1-score	support
0.0	0.90	0.98	0.94	1997
1.0	0.52	0.15	0.23	250
accuracy			0.89	2247
macro avg	0.71	0.57	0.59	2247
weighted avg	0.86	0.89	0.86	2247

### Accuracy, Precision, Recall, F1 Score

```
In [51]: acc_knn = accuracy_score(Y_test, Y_pred_knn_test)
prec_knn, rec_knn, f1_knn, sup_knn = precision_recall_fscore_support(Y_test, Y_pred_knn_test)
print('Accuracy Score =', acc_knn)
print('Precision =', prec_knn[0])
print('Recall =', rec_knn[0])
print('F1 Score =', f1_knn[0])
```

Accuracy Score = 0.8900756564307967  
 Precision = 0.9021139705882353  
 Recall = 0.9829744616925388  
 F1 Score = 0.940809968847352

### Checking for Overfitting and Underfitting

```
In [52]: knn_train_acc = accuracy_score(Y_train, Y_pred_knn_train)
```

```
knn_test_acc = accuracy_score(Y_test, Y_pred_knn_test)
print('Training Accuracy =', knn_train_acc)
print('Testing Accuracy =', knn_test_acc)
```

Training Accuracy = 0.9052760463045414  
Testing Accuracy = 0.8900756564307967

**There is no big variation in the training and testing accuracy. Therefore, the KNN Classifier model is not overfit or underfit.**

### Confusion Matrix

```
In [53]: pd.crosstab(Y_test.ravel(), Y_pred_knn_test)
```

```
Out[53]:
```

	col_0	0.0	1.0
row_0			
0.0	1963	34	
1.0	213	37	

## Random Forest Classifier

### Classification Report

```
In [54]: print(classification_report(Y_test, Y_pred_rf_test))
```

	precision	recall	f1-score	support
0.0	0.89	1.00	0.94	1997
1.0	0.00	0.00	0.00	250
accuracy			0.89	2247
macro avg	0.44	0.50	0.47	2247
weighted avg	0.79	0.89	0.84	2247

```
/opt/conda/envs/Python-3.9/lib/python3.9/site-packages/sklearn/metrics/_classification.py:
1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 i
n labels with no predicted samples. Use `zero_division` parameter to control this behavio
r.
    _warn_prf(average, modifier, msg_start, len(result))
/opt/conda/envs/Python-3.9/lib/python3.9/site-packages/sklearn/metrics/_classification.py:
1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 i
n labels with no predicted samples. Use `zero_division` parameter to control this behavio
r.
    _warn_prf(average, modifier, msg_start, len(result))
/opt/conda/envs/Python-3.9/lib/python3.9/site-packages/sklearn/metrics/_classification.py:
1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 i
n labels with no predicted samples. Use `zero_division` parameter to control this behavio
r.
    _warn_prf(average, modifier, msg_start, len(result))
```

### Accuracy, Precision, Recall, F1 Score

```
In [55]: acc_rf = accuracy_score(Y_test, Y_pred_rf_test)
prec_rf, rec_rf, f1_rf, sup_rf = precision_recall_fscore_support(Y_test, Y_pred_rf_test)
print('Accuracy Score =', acc_rf)
print('Precision =', prec_rf[0])
print('Recall =', rec_rf[0])
print('F1 Score =', f1_rf[0])
```

```
Accuracy Score = 0.8887405429461505
Precision = 0.8887405429461505
Recall = 1.0
F1 Score = 0.9410933081998115
```

```
/opt/conda/envs/Python-3.9/lib/python3.9/site-packages/sklearn/metrics/_classification.py:
1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in
labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
```

### Checking for Overfitting and Underfitting

```
In [56]: rf_train_acc = accuracy_score(Y_train, Y_pred_rf_train)
rf_test_acc = accuracy_score(Y_test, Y_pred_rf_test)
print('Training Accuracy =', rf_train_acc)
print('Testing Accuracy =', rf_test_acc)
```

```
Training Accuracy = 0.8756678539626002
Testing Accuracy = 0.8887405429461505
```

**There is no big variation in the training and testing accuracy. Therefore, the Random Forest Classifier model is not overfit or underfit.**

### Confusion Matrix

```
In [57]: pd.crosstab(Y_test.ravel(), Y_pred_rf_test)
```

```
Out[57]:
```

	col_0	0.0
row_0	0.0	1997
1.0	250	

*On comparing the four models built, based on the performance metrics it is clear that Support Vector Classifier gives the highest performance. Hence, that model is chosen for deployment*

## Dumping the Chosen Model into pkl file

```
In [58]: joblib.dump(svc, 'model.pkl')
```

```
Out[58]: ['model.pkl']
```

## Dumping the Scaling Modules into pkl file

```
In [59]: pickle.dump(scale_crs, open('crs_scale.pkl', 'wb'))
pickle.dump(scale_flnum, open('flnum_scale.pkl', 'wb'))
```

```
In [60]: !pip install -U ibm-watson-machine-learning
```

Requirement already satisfied: ibm-watson-machine-learning in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (1.0.257)

Requirement already satisfied: pandas<1.5.0,>=0.24.2 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-watson-machine-learning) (1.3.4)

Requirement already satisfied: lomond in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-watson-machine-learning) (0.3.3)

Requirement already satisfied: tabulate in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-watson-machine-learning) (0.8.9)

Requirement already satisfied: ibm-cos-sdk==2.11.\* in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-watson-machine-learning) (2.11.0)

Requirement already satisfied: urllib3 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-watson-machine-learning) (1.26.7)

Requirement already satisfied: packaging in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-watson-machine-learning) (21.3)

Requirement already satisfied: certifi in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-watson-machine-learning) (2022.9.24)

Requirement already satisfied: importlib-metadata in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-watson-machine-learning) (4.8.2)

Requirement already satisfied: requests in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-watson-machine-learning) (2.26.0)

Requirement already satisfied: ibm-cos-sdk-core==2.11.0 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-cos-sdk==2.11.\*->ibm-watson-machine-learning) (2.11.0)

Requirement already satisfied: jmespath<1.0.0,>=0.7.1 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-cos-sdk==2.11.\*->ibm-watson-machine-learning) (0.10.0)

Requirement already satisfied: ibm-cos-sdk-s3transfer==2.11.0 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-cos-sdk==2.11.\*->ibm-watson-machine-learning) (2.11.0)

Requirement already satisfied: python-dateutil<3.0.0,>=2.1 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from ibm-cos-sdk-core==2.11.0->ibm-cos-sdk==2.11.\*->ibm-watson-machine-learning) (2.8.2)

Requirement already satisfied: pytz>=2017.3 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from pandas<1.5.0,>=0.24.2->ibm-watson-machine-learning) (2021.3)

Requirement already satisfied: numpy>=1.17.3 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from pandas<1.5.0,>=0.24.2->ibm-watson-machine-learning) (1.20.3)

Requirement already satisfied: six>=1.5 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from python-dateutil<3.0.0,>=2.1->ibm-cos-sdk-core==2.11.0->ibm-cos-sdk==2.11.\*->ibm-watson-machine-learning) (1.15.0)

Requirement already satisfied: idna<4,>=2.5 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from requests->ibm-watson-machine-learning) (3.3)

Requirement already satisfied: charset-normalizer~=2.0.0 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from requests->ibm-watson-machine-learning) (2.0.4)

Requirement already satisfied: zipp>=0.5 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from importlib-metadata->ibm-watson-machine-learning) (3.6.0)

Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /opt/conda/envs/Python-3.9/lib/python3.9/site-packages (from packaging->ibm-watson-machine-learning) (3.0.4)

```
In [61]: from ibm_watson_machine_learning import APIClient
import json
```

```
In [62]: wml_credentials = {
    "apikey": "G0-fTz5hk2xE6cxkk-PyyzQmio8NNqeMG_wTLTgCSqmU",
    "url": "https://us-south.ml.cloud.ibm.com"
}
```

```
In [63]: wml_clients = APIClient(wml_credentials)
```

```
In [64]: wml_clients.spaces.list()
```

Note: 'limit' is not provided. Only first 50 records will be displayed if the number of records exceed 50

ID	NAME	CREATED
267b8ab3-7a6f-4efa-81fb-bc602a9f29e2	flight	2022-11-19T09:48:15.358Z



```
In [65]: space_id = "267b8ab3-7a6f-4efa-81fb-bc602a9f29e2"
```

```
In [66]: wml_clients.set.default_space(space_id)
```

```
Out[66]: 'SUCCESS'
```

```
In [67]: wml_clients.software_specifications.list(500)
```

NAME	ASSET_ID	TYPE
default_py3.6	0062b8c9-8b7d-44a0-a9b9-46c416adcbd9	base
kernel-spark3.2-scala2.12	020d69ce-7ac1-5e68-ac1a-31189867356a	base
pytorch-onnx_1.3-py3.7-edt	069ea134-3346-5748-b513-49120e15d288	base
scikit-learn_0.20-py3.6	09c5a1d0-9c1e-4473-a344-eb7b665ff687	base
spark-mllib_3.0-scala_2.12	09f4cff0-90a7-5899-b9ed-1ef348aebdee	base
pytorch-onnx_rt22.1-py3.9	0b848dd4-e681-5599-be41-b5f6fccc6471	base
ai-function_0.1-py3.6	0cdb0f1e-5376-4f4d-92dd-da3b69aa9bda	base
shiny-r3.6	0e6e79df-875e-4f24-8ae9-62dcc2148306	base
tensorflow_2.4-py3.7-horovod	1092590a-307d-563d-9b62-4eb7d64b3f22	base
pytorch_1.1-py3.6	10ac12d6-6b30-4ccd-8392-3e922c096a92	base
tensorflow_1.15-py3.6-ddl	111e41b3-de2d-5422-a4d6-bf776828c4b7	base
autoai-kb_rt22.2-py3.10	125b6d9a-5b1f-5e8d-972a-b251688ccf40	base
runtime-22.1-py3.9	12b83a17-24d8-5082-900f-0ab31fbfd3cb	base
scikit-learn_0.22-py3.6	154010fa-5b3b-4ac1-82af-4d5ee5abbc85	base
default_r3.6	1b70aec3-ab34-4b87-8aa0-a4a3c8296a36	base
pytorch-onnx_1.3-py3.6	1bc6029a-cc97-56da-b8e0-39c3880dbbe7	base
kernel-spark3.3-r3.6	1c9e5454-f216-59dd-a20e-474a5cdf5988	base
pytorch-onnx_rt22.1-py3.9-edt	1d362186-7ad5-5b59-8b6c-9d0880bde37f	base
tensorflow_2.1-py3.6	1eb25b84-d6ed-5dde-b6a5-3fbdf1665666	base
spark-mllib_3.2	20047f72-0a98-58c7-9ff5-a77b012eb8f5	base
tensorflow_2.4-py3.8-horovod	217c16f6-178f-56bf-824a-b19f20564c49	base
runtime-22.1-py3.9-cuda	26215f05-08c3-5a41-a1b0-da66306ce658	base
do_py3.8	295addb5-9ef9-547e-9bf4-92ae3563e720	base
autoai-ts_3.8-py3.8	2aa0c932-798f-5ae9-abd6-15e0c2402fb5	base
tensorflow_1.15-py3.6	2b73a275-7cbf-420b-a912-eae7f436e0bc	base
kernel-spark3.3-py3.9	2b7961e2-e3b1-5a8c-a491-482c8368839a	base
pytorch_1.2-py3.6	2c8ef57d-2687-4b7d-acce-01f94976dac1	base
spark-mllib_2.3	2e51f700-bca0-4b0d-88dc-5c6791338875	base
pytorch-onnx_1.1-py3.6-edt	32983cea-3f32-4400-8965-dde874a8d67e	base
spark-mllib_3.0-py37	36507ebe-8770-55ba-ab2a-eafe787600e9	base
spark-mllib_2.4	390d21f8-e58b-4fac-9c55-d7ceda621326	base
autoai-ts_rt22.2-py3.10	396b2e83-0953-5b86-9a55-7ce1628a406f	base
xgboost_0.82-py3.6	39e31acd-5f30-41dc-ae44-60233c80306e	base
pytorch-onnx_1.2-py3.6-edt	40589d0e-7019-4e28-8daa-fb03b6f4fe12	base
pytorch-onnx_rt22.2-py3.10	40e73f55-783a-5535-b3fa-0c8b94291431	base
default_r36py38	41c247d3-45f8-5a71-b065-8580229facf0	base
autoai-ts_rt22.1-py3.9	4269d26e-07ba-5d40-8f66-2d495b0c71f7	base
autoai-obm_3.0	42b92e18-d9ab-567f-988a-4240ba1ed5f7	base
pmml-3.0_4.3	493bcb95-16f1-5bc5-bee8-81b8af80e9c7	base
spark-mllib_2.4-r_3.6	49403dff-92e9-4c87-a3d7-a42d0021c095	base
xgboost_0.90-py3.6	4ff8d6c2-1343-4c18-85e1-689c965304d3	base
pytorch-onnx_1.1-py3.6	50f95b2a-bc16-43bb-bc94-b0bed208c60b	base
autoai-ts_3.9-py3.8	52c57136-80fa-572e-8728-a5e7cbb42cde	base
spark-mllib_2.4-scala_2.11	55a70f99-7320-4be5-9fb9-9edb5a443af5	base
spark-mllib_3.0	5c1b0ca2-4977-5c2e-9439-ffd44ea8ffe9	base
autoai-obm_2.0	5c2e37fa-80b8-5e77-840f-d912469614ee	base
spss-modeler_18.1	5c3cad7e-507f-4b2a-a9a3-ab53a21dee8b	base
cuda-py3.8	5d3232bf-c86b-5df4-a2cd-7bb870a1cd4e	base
runtime-22.2-py3.10-xc	5e8cddff-db4a-5a6a-b8aa-2d4af9864dab	base
autoai-kb_3.1-py3.7	632d4b22-10aa-5180-88f0-f52dfb6444d7	base
pytorch-onnx_1.7-py3.8	634d3cdc-b562-5bf9-a2d4-ea90a478456b	base
spark-mllib_2.3-r_3.6	6586b9e3-ccd6-4f92-900f-0f8cb2bd6f0c	base
tensorflow_2.4-py3.7	65e171d7-72d1-55d9-8ebb-f813d620c9bb	base
spss-modeler_18.2	687eddc9-028a-4117-b9dd-e57b36f1efa5	base
pytorch-onnx_1.2-py3.6	692a6a4d-2c4d-45ff-a1ed-b167ee55469a	base
spark-mllib_2.3-scala_2.11	7963efe5-bbec-417e-92cf-0574e21b4e8d	base
spark-mllib_2.4-py37	7abc992b-b685-532b-a122-a396a3cdbaab	base
caffe_1.0-py3.6	7bb3dbe2-da6e-4145-918d-b6d84aa93b6b	base
pytorch-onnx_1.7-py3.7	812c6631-42b7-5613-982b-02098e6c909c	base
cuda-py3.6	82c79ece-4d12-40e6-8787-a7b9e0f62770	base
tensorflow_1.15-py3.6-horovod	8964680e-d5e4-5bb8-919b-8342c6c0dfd8	base
hybrid_0.1	8c1a58c6-62b5-4dc4-987a-df751c2756b6	base
pytorch-onnx_1.3-py3.7	8d5d8a87-a912-54cf-81ec-3914adaa988d	base
caffe-ibm_1.0-py3.6	8d863266-7927-4d1e-97d7-56a7f4c0a19b	base

runtime-22.2-py3.10-cuda	8ef391e4-ef58-5d46-b078-a82c211c1058	base
spss-modeler_17.1	902d0051-84bd-4af6-ab6b-8f6aa6fdeabb	base
do_12.10	9100fd72-8159-4eb9-8a0b-a87e12eeefa36	base
do_py3.7	9447fa8b-2051-4d24-9eef-5acb0e3c59f8	base
spark-mllib_3.0-r_3.6	94bb6052-c837-589d-83f1-f4142f219e32	base
cuda-py3.7-opence	94e9652b-7f2d-59d5-ba5a-23a414ea488f	base
nlp-py3.8	96e60351-99d4-5a1c-9cc0-473ac1b5a864	base
cuda-py3.7	9a44990c-1aa1-4c7d-baf8-c4099011741c	base
hybrid_0.2	9b3f9040-9cee-4ead-8d7a-780600f542f7	base
spark-mllib_3.0-py38	9f7a8fc1-4d3c-5e65-ab90-41fa8de2d418	base
autoai-kb_3.3-py3.7	a545cca3-02df-5c61-9e88-998b09dc79af	base
spark-mllib_3.0-py39	a6082a27-5acc-5163-b02c-6b96916eb5e0	base
runtime-22.1-py3.9-do	a7e7dbf1-1d03-5544-994d-e5ec845ce99a	base
default_py3.8	ab9e1b80-f2ce-592c-a7d2-4f2344f77194	base
tensorflow_rt22.1-py3.9	acd9c798-6974-5d2f-a657-ce06e986df4d	base
kernel-spark3.2-py3.9	ad7033ee-794e-58cf-812e-a95f4b64b207	base
autoai-obm_2.0 with Spark 3.0	af10f35f-69fa-5d66-9bf5-acb58434263a	base
runtime-22.2-py3.10	b56101f1-309d-549b-a849-eea63f77b2fb	base
default_py3.7_opence	c2057dd4-f42c-5f77-a02f-72bdbd3282c9	base
tensorflow_2.1-py3.7	c4032338-2a40-500a-beef-b01ab2667e27	base
do_py3.7_opence	cc8f8976-b74a-551a-bb66-6377f8d865b4	base
spark-mllib_3.3	d11f2434-4fc7-58b7-8a62-755da64fdaf8	base
autoai-kb_3.0-py3.6	d139f196-e04b-5d8b-9140-9a10ca1fa91a	base
spark-mllib_3.0-py36	d82546d5-dd78-5fbb-9131-2ec309bc56ed	base
autoai-kb_3.4-py3.8	da9b39c3-758c-5a4f-9cfd-457dd4d8c395	base
kernel-spark3.2-r3.6	db2fe4d6-d641-5d05-9972-73c654c60e0a	base
autoai-kb_rt22.1-py3.9	db6afe93-665f-5910-b117-d879897404d9	base
tensorflow_rt22.1-py3.9-horovod	dda170cc-ca67-5da7-9b7a-cf84c6987fae	base
autoai-ts_1.0-py3.7	deef04f0-0c42-5147-9711-89f9904299db	base
tensorflow_2.1-py3.7-horovod	e384fce5-fdd1-53f8-bc71-11326c9c635f	base
default_py3.7	e4429883-c883-42b6-87a8-f419d64088cd	base
do_22.1	e51999ba-6452-5f1f-8287-17228b88b652	base
autoai-obm_3.2	eae86aab-da30-5229-a6a6-1d0d4e368983	base
runtime-22.2-r4.2	ec0a3d28-08f7-556c-9674-ca7c2dba30bd	base
tensorflow_rt22.2-py3.10	f65bd165-f057-55de-b5cb-f97cf2c0f393	base
do_20.1	f686cdd9-7904-5f9d-a732-01b0d6b10dc5	base
pytorch-onnx_rt22.2-py3.10-edt	f8a05d07-e7cd-57bb-a10b-23f1d4b837ac	base
scikit-learn_0.19-py3.6	f963fa9d-4bb7-5652-9c5d-8d9289ef6ad9	base
tensorflow_2.4-py3.8	fe185c44-9a99-5425-986b-59bd1d2eda46	base
-----	-----	----

```
In [68]: import sklearn
sklearn.__version__
'1.0.2'
```

```
Out[68]: '1.0.2'
```

```
In [69]: MODEL_NAME="supportvectormachine"
DEPLOYMENT_NAME="svc_deployment"
DEMO_MODEL=svc
```

```
In [70]: soft_sepc_id=wml_clients.software_specifications.get_id_by_name("runtime-22.1-py3.9")
```

```
In [71]: model_props={
    wml_clients.repository.ModelMetaNames.NAME:MODEL_NAME,
    wml_clients.repository.ModelMetaNames.TYPE:"scikit-learn_1.0",
    wml_clients.repository.ModelMetaNames.SOFTWARE_SPEC_UID: soft_sepc_id
}
```

```
In [72]: model_details=wml_clients.repository.store_model(model=DEMO_MODEL,meta_props=model_props,t
    training_target=Y_train.ravel())
```

```
In [73]: model_details
```

```
Out[73]: {'entity': {'hybrid_pipeline_software_specs': [],
  'label_column': 'l1',
  'schemas': {'input': [{'fields': [{'name': 'f0', 'type': 'float'},
    {'name': 'f1', 'type': 'float'},
    {'name': 'f2', 'type': 'float'},
    {'name': 'f3', 'type': 'float'},
    {'name': 'f4', 'type': 'float'},
    {'name': 'f5', 'type': 'float'},
    {'name': 'f6', 'type': 'float'},
    {'name': 'f7', 'type': 'float'},
    {'name': 'f8', 'type': 'float'},
    {'name': 'f9', 'type': 'float'},
    {'name': 'f10', 'type': 'float'},
    {'name': 'f11', 'type': 'float'},
    {'name': 'f12', 'type': 'float'},
    {'name': 'f13', 'type': 'float'},
    {'name': 'f14', 'type': 'float'},
    {'name': 'f15', 'type': 'float'}]},
  'id': '1',
  'type': 'struct'}],
  'output': []},
  'software_spec': {'id': '12b83a17-24d8-5082-900f-0ab31fbfd3cb',
  'name': 'runtime-22.1-py3.9'},
  'type': 'scikit-learn_1.0'},
  'metadata': {'created_at': '2022-11-19T19:28:04.765Z',
  'id': 'e3e793cc-079e-4363-9775-0a67bb3665b3',
  'modified_at': '2022-11-19T19:28:07.620Z',
  'name': 'supportvectormachine',
  'owner': 'IBMid-661003XLH3',
  'resource_key': '2a8d6e41-3599-4d66-aa9d-aaa87b678c20',
  'space_id': '267b8ab3-7a6f-4efa-81fb-bc602a9f29e2'},
  'system': {'warnings': []}}
```

```
In [74]: model_id = wml_clients.repository.get_model_id(model_details)
model_id
```

```
Out[74]: 'e3e793cc-079e-4363-9775-0a67bb3665b3'
```

```
In [75]: deployment_props = {
  wml_clients.deployments.ConfigurationMetaNames.NAME: DEPLOYMENT_NAME,
  wml_clients.deployments.ConfigurationMetaNames.ONLINE: {}
}
```

```
In [76]: deployment = wml_clients.deployments.create(
  artifact_uid=model_id,
  meta_props=deployment_props
)
```

```
#####
```

```
Synchronous deployment creation for uid: 'e3e793cc-079e-4363-9775-0a67bb3665b3' started
```

```
#####
```

```
initializing
```

```
Note: online_url is deprecated and will be removed in a future release. Use serving_urls instead.
```

```
ready
```

```
-----  
-----  
Successfully finished deployment creation, deployment_uid='486b7d6a-343e-4048-bee1-e73454d27219'  
-----  
-----
```

```
In [ ]:
```